

Analysis of the Quantum Advantages for Deep Hedging

Soham Deshpande

Supervisor: Dr. Srinandan Dasmahapatra

Second Supervisor: Dr. Hansung Kim

April 2025

A project report submitted for the award of MEng Computer Science

Abstract

Parameterised Quantum Circuits (PQCs) have opened many doors, one such being the use in financial markets. In this paper, I look at the problem of developing an accurate market generator through the use of quantum computing for the purposes of hedging. Given a Quantum Circuit Born Machine (QCBM), we are able to exploit the high expressibility to generate synthetic data that mimics the statistical distribution of the original dataset. The market generator is then used to simulate an underlying asset to maturity with the intent of learning an optimal hedging strategy π^* , a showcase of a data-driven approach to hedging exposure. I show that the synthetic data produced by this method has shown to capture the fat tails of the market better than classical methods, as well as demonstrating superiority in out-of-sample testing with COVID data. Different generator architectures have been compared to maximise the quality of the synthetic data and avoid issues with barren plateaus. The findings of this research will contribute to the growing literature on risk management in quantitative finance, with applications of the market generator extending beyond deep hedging.

Contents

1	Problem Statement	1
2	Related Literature	2
3	Markets and Derivatives	4
3.1	Brownian Motion	4
3.1.1	Itô Process	4
3.1.2	Geometric Brownian Motion	4
3.2	Market	5
3.3	Derivatives	5
3.3.1	Futures	5
3.3.2	Options	6
3.4	Market Data	6
3.4.1	Euro Stoxx 50	6
3.4.2	Brent Crude Oil	6
4	Merton-Jump Diffusion Model	8
4.1	Black-Scholes	8
4.2	Model	8
4.3	Assumptions & Limitations	9
4.4	Calibration	9
5	Quantum Computing	11
5.1	Quantum Systems	11
5.2	Born Rule	12
5.3	State Preparation	12
5.4	Parameterised Quantum Circuits	12
6	Quantum Circuit Born Machine	13
6.1	Barren Plateau	13
6.2	Architectures	13
6.2.1	Brick	14
6.2.2	Pyramid	14
6.2.3	Butterfly	14
6.3	ZX Calculus	14
7	Results	15
7.1	Path Generation	15
7.2	Volatility Analysis	18
7.3	VaR & CVaR	21
7.4	Hedging	24
7.5	Barren Plateau	24
7.5.1	ZX Calculus	24
8	Outlook and Conclusions	25

9	Appendix	26
9.1	QCBM Architectures	26

1 Problem Statement

The problem of hedging a portfolio of derivatives is an important part of risk management used widely in financial institutions. This involves understanding the exposure to the market, and taking strategic positions to negate some of this risk. In an ideal world we can picture a perfect, frictionless market where transaction costs are negligible and every asset in the space has a price; here we can price and hedge perfectly. Unfortunately in practice, we experience incomplete markets due to frictions, costs that interfere with trades such as transaction costs or imperfect information. In addition, recent years have presented markets with periods of heightened volatility, much that disobey traditional frameworks. This generates the need for complex, realistic market models that can account for these.

Traditional methods of hedging options has shown to be ineffective for equity markets, new information resulting in rapid changes. Much of the available literature models the market as a smooth, continuous stochastic process within a Gaussian space. Such models are sufficient for common market activity but fail when presented with discontinuous moves in price. These can be reactions to geopolitical events or natural disasters; traditional models are incapable of capturing the effects. The introduction of Jump-diffusion models aimed to solve this issue though face similar issues. In reaction, we have recently observed non-parametric models that harness neural networks and machine learning which aim to demonstrate high accuracy on out-of-sample forecasts.

An alternative approach that has recently emerged utilises the power of quantum computing. The introduction of parameterised quantum circuits(PQCs) have opened up new pathways for navigating complex, large scale time series datasets. Through rotation gates and quantum entanglement, we are able to learn complex distributions and relationships.

In this research, I aim to tackle the problem of generating synthetic financial data, addressing issues that come about from using a classical method, particularly the estimation of tail risk and skewness. Through comparisons between traditional approaches, I aim to demonstrate an advantage in the expressibility of Quantum Circuit Born Machines (QCBMs); these will be described quantitatively using measures such as Value at Risk (VaR) and Conditional Value at Risk (CVaR). By performing out-of-sample tests on COVID and Oil stock price data, I aim to highlight the weaknesses of traditional models and showcase a quantum superiority. There will also be an exploration into the variety of architectures available for the QCBM, evaluating different ansatz designs. Where unusual behaviours due to the quantum nature occur, such as barren plateau, I will explore in greater detail as well as any circuit optimisation techniques that may present themselves as possible solutions.

This paper will aim to add to the existing literature on risk management for financial firms as well as providing a framework for generating synthetic data. In addition to that, I aim to extend to the QCBM research that exists currently, noting down any behaviours that may be of interest to the curious and potential experts in the field.

2 Related Literature

To place this research within the context of existing literature, we can split the project into 2 components: the market generator, and parameterised quantum circuits.

The work around deep hedging has evolved, moving away from Greek-based hedging towards a sturdier framework using machine learning. Here a lot of work is being done, with many papers emphasising on using neural networks for optimising delta and gamma exposure [1, 33]. Buehler introduced an approach, modelling trading decisions as neural networks instead of relying on parameterised models [8]. Subsequent advancements focussed on developing realistic market simulators. Wissel proposed a market model for path generation of options but this still employed risk-neutral diffusion[38]. Wiese then introduced a new dimension by using GANs to convert options into local volatility models with simpler no-arbitrage constraints. This focussed on the local stochastic nature of options [10, 45, 46]. Some approaches suggest using actor-critic reinforcement learning algorithms to solve for an optimal value function, a move towards searching for a global maximum over local risk management [7, 27].

Recent research explores using quantum computing to hedge portfolios, here the authors presented a quantum reinforcement learning method based on policy-search and distributional actor-critic algorithms. They proposed using a Quantum Neural Network to approximate the value of a given value function by predicting the expected utility of returns using compound and orthogonal layers which were built using Hamming-weight unitaries [23].

TO CHANGE: This helped overcome the barren plateau by ensuring the gradient variance does not vanish exponentially with qubit count.

Another method models the entire return distribution, leveraging parameterised circuits to learn categorical distributions and capture variability and tail risk [9, 12].

There is an immense amount of research being done on exploiting the benefits of quantum computing, recent advancements being in quantum algorithms. These claim to provide exponential speed-up over classical methods, though in reality, we see great complexity in state preparation, requiring $\Theta(2^n/n)$ circuit depth with n qubits or $\Theta(n)$ circuit depth with $\Theta(2^n)$ ancillary qubits[48]. Here we see hybrid models such as Born machines and Quantum generative adversarial networks boasting high generalisation ability [18, 20, 22].

There has also been research in harnessing back action from quantum weak measurements to enhance the ability of quantum machine learning algorithms. In quantum reservoir computing, the ability to retain information from past inputs plays a key role in processing temporal series and producing future predictions [16, 17, 19, 28].

This research aims to combine the needs of financial firms in hedging portfolios using realistic market models by utilising QCBMs as a tool for simulating paths in combination with deep hedging engines for learning optimal policies. A comparison will be made

against hedging under Merton-Jump diffusion.

3 Markets and Derivatives

The market, though inherently can be thought of as a completely random process, where bids and asks are fulfilled, can be modelled as a stochastic process. The aim of this chapter is to serve as a brief introduction and set up notation for later chapters.

3.1 Brownian Motion

To represent this stochasticity, we must employ techniques introduced by Norbert Wiener, the Wiener process, more commonly referred to as standard Brownian Motion. This framework allows us to model continuous random walks of our stock price. Formally, a standard Wiener process, W_t , is a stochastic process where

1. $W_0 = 0$
2. The process W_t has stationary, independent increments
3. $\forall t \in \mathbb{Z}$, the random variable W_t is normally distributed, $N(0, t)$
4. The paths of W_t are continuous ensuring no jumps in the path trajectory

These assumptions will help us understand the shortfalls of traditional techniques.

3.1.1 Itô Process

Itô processes are crucial for understanding the mathematical set up for modelling our assets. Itô calculus allows us to extend our understanding of deterministic calculus to the realm of stochasticity. Suppose X_t Itô process can be defined as a stochastic process which can be written in the form

$$X_t = X_0 + \int_0^t U_s ds + \int_0^t V_s dW_s \quad (3.1)$$

3.1.2 Geometric Brownian Motion

We can extend Brownian Motion to Geometric Brownian Motion by exponentiating the BM; this is done to satisfy the condition that stock prices are non-negative. We can now consider a continuous time process $S(t)$ which satisfies the SDE

$$dS_t = \mu S(t)dt + \sigma S(t)dW_t \quad (3.2)$$

where μ is the drift parameter, σ is the volatility parameter, and W_t is a Wiener process. The solution is given by

$$S_t = S_0 \exp \left[\left(\mu - \frac{1}{2}\sigma^2 \right)t + \sigma W_t \right] \quad (3.3)$$

3.2 Market

Consider a market with a finite time horizon T defined on the probability space (Ω, \mathcal{F}, P) along with a filtration $\mathbf{F} = \{\mathcal{F}|0 \leq t \leq T\}$. This can be thought of as an adapted $(n+1)$ dimensional Itô process $X(t) = (X_0(t), X_1(t), \dots, X_n(t))$ which has the form

$$dX_0(t) = \rho(t, \omega)X_0(t)dt; \quad X_0(0) = 1 \quad (3.4)$$

and

$$dX_i = \mu_i(t, \omega)dt + \sigma_i(t, \omega)dB(t); \quad X_i(0) = x_i \quad (3.5)$$

where $X_i(t)$ is the price of asset i at a given time t .

We can define a portfolio in the market as

$$\theta(t, \omega) = (\theta_0(t, \omega), \theta_1(t, \omega), \dots, \theta_n(t, \omega)) \quad (3.6)$$

where the components $\theta_n(t, \omega)$ represents the number of units of a given asset held at time t .

Following from that, we can define the value of a given portfolio to be

$$V(t, \omega) = V^\theta(t, \omega) = \sum_{i=0}^n \theta_i(t)X_i(t) \quad (3.7)$$

Lastly, it is important to state that the portfolio is self-financing, any trading strategy α requires no extra cost beyond the initial capital

$$V(t) = V(0) + \int_0^t \theta(s) \cdot dX(s) \quad (3.8)$$

We can also make the following assumptions about the market:

- The market is liquid, allowing the trade to execute instantaneously
- There is no bid-ask spread, the price to buy and sell is the same
- Trading actions taken have no impact on the price of the asset traded

3.3 Derivatives

A derivative refers to any financial instrument whose value is derived from an underlying security, the most fundamental being futures and options. It is common practice to refer to the given underlying security as just 'underlying'.

3.3.1 Futures

A futures contract is a contract that gives the right and obligation to buy a given asset i at a specified time T at price K .

3.3.2 Options

The two types of options going to be explored are Puts and Calls; a Call option gives the owner the right but not the obligation to buy a given asset i at a specified price K at time T . Similar to the Call, a Put option gives the owner the right but not the obligation to sell a given asset i at a price K at time T . If the owner can exercise the option any time up to T , we call this an American option. For the purposes of this research, I will only be dealing with vanilla European options.

It is important to define the payoffs for both options:

$$C_T = \max(0, S_T - K) \quad (3.9)$$

$$P_T = \max(0, K - S_T) \quad (3.10)$$

3.4 Market Data

In this research I will be focussing on hedging a portfolio consisting of a single asset, hence requiring a simulation of a single underlying.

3.4.1 Euro Stoxx 50

The Euro Stoxx 50 Index (SX5E) and relevant derivatives. This is a stock index of 50 stocks in the Eurozone. This index captures around 60% of the free-float market capitalisation of the Euro Stoxx Total Market Index which covers about 95% of the free-float market in the Eurozone[14]. Rationale behind choosing this index is the availability of data, options traded with SX5E as the underlying and the liquidity of the index.

Derivatives that are held in the portfolio to be hedged will include those that have SX5E as the underlying, examples are weekly, monthly, and quarterly expiration options. These are European-style so can only be exercised upon maturity. Data can be found on Bloomberg[5] and Refinitiv [26].




Figure 1: Euro Stoxx 50 price chart

3.4.2 Brent Crude Oil

Commodities often have complex dynamics, driven by geopolitical events and natural supply and demand. As well as this, oil in particular lends itself to increased volatility often responding to changes by the Federal Reserve and OPEC(Organisation of petroleum-exporting countries). These factors create characteristics such as heavy tails and jumps,

often not being represented well by traditional models. I will be using Brent Oil as a benchmark asset to compare the expressibility of quantum derived models vs existing classical ones.

4 Merton-Jump Diffusion Model

My model of choice for comparison is the Merton-Jump Diffusion model, this an elementary model that goes beyond Black-Scholes by trying to capture the negative skewness and excess kurtosis of log price returns. This is done through the addition of a compound Poisson jump process. This aims to represent the jumps we observe in the market in a more realistic fashion rather than assuming constant volatility assumption made by Black-Scholes. As a simplification, I will be referring to Black-Scholes by BS and Merton-Jump Diffusion with MJD.

4.1 Black-Scholes

Let's start with the Black-Scholes model, an elementary model first proposed in (give year here) to price European vanilla options.

Consider a call options on a non-dividend paying stock with expiry T and strike K . We will assume that the asset price will obey geometric brownian motion hence giving:

$$dS_t = \mu S_t dt + \sigma S_t dW_t \quad (4.1)$$

where W_t is a standard Brownian motion. We assume interest rates to be constant, meaning a unit of a given currency at time t will be worth e^{rt} at time t .

We can now consider the value of our call option, C , at time t , by way of Itô's lemma we can say that

$$dC = () \quad (4.2)$$

4.2 Model

A standard derivation of the model will allow us to explore its assumptions and limitations. This model consists of two components, jump and diffusion. The diffusion will be modelled using a Weiner process and log-normal jumps driven by a Poisson process. This gives us the following SDE.

$$dS_t = (\alpha - \lambda k)S_t dt + \sigma S_t dW_t + (y_t - 1)S_t dN_t \quad (4.3)$$

where W_t and N_t are Weiner and Poisson processes respectively. λ represents the intensity of the jumps, α is the drift rate (expected return), and k is the expected jump size. Solving the equation gives us an exponential Lévy model described by

$$S_t = S_0 e^{\mathcal{L}_t} \quad (4.4)$$

where S_t is the stock price at time t , S_0 is the initial stock price. We can also define \mathcal{L}_t to be

$$\mathcal{L}_t = \left(\alpha - \frac{\sigma^2}{2} - \lambda k\right)t + \sigma W_t + \sum_{i=1}^{N_t} Y_i \quad (4.5)$$

4.3 Assumptions & Limitations

Through inspection of the equations, we can observe the following assumptions:

1. The asset price experiences continuous, random fluctuations over time, governed by Brownian motion (GBM)
2. The asset price experiences sudden, discontinuous jumps modelled by a Poisson process, occurring at a constant rate λ
3. Jumps sizes are assumed to be log-normal $\ln(y_t) \sim \mathcal{N}(\mu, \sigma^2)$

Starting with the first assumption, we can see that assuming GBM may produce unrealistic behaviour, most important being a lack of excess kurtosis. Markets often exhibit fat tails, especially within commodities. One such event may be the release of news from OPEC+, the organisation of petroleum-exporting countries. A restriction in oil production may cause the price of oil to jump rapidly. In recent times, wars and conflict has also become ever present, causing large movements in asset prices; therefore it is not unrealistic to expect extreme price movements to be more frequent than can be modelled by a Gaussian.

The MJD requires calibration of parameters before use, typically done using historical data or implied volatility surfaces. Once calibrated these become assumptions of the data and so do not change even if the market observations move away from it. This would lead us to expect higher overfitting to the historical data, possibly failing in unseen conditions such as the market's reaction to COVID.

We also may expect poor volatility clustering with the MJD; constant volatility is not experienced by the market, instead periods of high volatility followed by periods of low volatility is observed. Though this paper won't be focussing on this phenomenon, it is important to consider.

4.4 Calibration

In this research, I have chosen to use maximum likelihood estimation to estimate the parameters for the MJD model. In the analytical solution we require five parameters: α , σ , μ_j , δ , and λ . These are the expected return, volatility of the given asset, expectation of the jump size, standard deviation of the jump size and lastly the jump intensity. We can then use MLE on the probability density of log returns $S_t = \ln(\frac{S_t}{S_0})$

$$P(S_t) = \sum_{i=0}^{\infty} \frac{e^{-\lambda t} (\lambda t)^i}{i!} N(S_t; (\alpha - \frac{\sigma^2}{2} - \lambda k)t + i\mu_j, \sigma^2 t + i\delta^2) \quad (4.6)$$

The likelihood function hence becomes

$$L(\theta; S) = \prod_{t=1}^T P(S_t) \quad (4.7)$$

We can minimise the negative log-likelihood to obtain

$$-\ln L(\theta; S) = -\sum_{t=1}^T \ln P(S_t) \quad (4.8)$$

Another popular option to calibrate the MJD model is by considering the implied volatility surface of existing options. This technique can lead to a calibration but suffers with issues surrounding the sensitivity of the tails of the asset prices. It is also well documented that given a function that measures the calibration error, we can observe a largely flat landscape surrounding the optimal solution, implying obtaining accurate parameters can become very computationally expensive, often requiring hundreds of iterations [jump05]. These difficulties can translate into a poor hedge, leaving a buyer overexposed to market fluctuations.

5 Quantum Computing

This section aims to serve as a brief introduction to quantum computing, this may prove to be pedestrian to the experienced but is included to maintain the accessibility of this research report.

5.1 Quantum Systems

Unlike classical computing, quantum computing acts in a non-deterministic manner, the computer remains in multiple states with given probabilities rather than a fixed resultant state as expected from classical computers. Formally we can define a qubit to be a quantum system where the states of 0 and 1 are represented by a pair of normalised and mutually orthogonal quantum states $|0\rangle$ and $|1\rangle$. Intuitively however, let's start with a two-state machine; we can describe such system to be in the state $|0\rangle$ with some amplitude α and in $|1\rangle$ with amplitude β . This can be represented as

$$|\psi\rangle = \alpha|0\rangle + \beta|1\rangle \quad (5.1)$$

for some α and β such that $|\alpha|^2 + |\beta|^2 = 1$; we can refer to this as a superposition. If measured in the standard basis, we would expect the outcome to be $|k\rangle$ with a certain probability, this outcome resulting in the output state of the measurement gate to also be $|k\rangle$. This would mean our state $|\psi\rangle$ is irreversibly lost; we refer to this as a collapse of state. Each qubit can be thought of as a vector, \mathbf{v} , on a Bloch's sphere which can be represented in two basis: θ and ψ . θ is the angle between \mathbf{v} and the z-axis. ψ becomes the angle between \mathbf{v} and the x-axis. Considering a more general parameterisation of 5.1 gives us

$$|\psi\rangle = \cos(\frac{\theta}{2})e^{i\psi_0}|0\rangle + \sin(\frac{\theta}{2})e^{i\psi_1}|1\rangle \quad (5.2)$$

TALK MORE HERE

When extending to a n-qubit system

Before forming quantum circuits, we must first understand how quantum gates operate. Quantum gates can be thought of as fixed unitary operations on selected qubits, often represented with a U . There are many gates but the ones we are concerned with for this project are the T , R_x , R_y , R_z and $CNOT$ gates. Given an angle θ we can form the universal gates:

$$R_x(\theta) = \begin{bmatrix} \cos(\theta/2) & -i \sin(\theta/2) \\ -i \sin(\theta/2) & \cos(\theta/2) \end{bmatrix}$$

$$R_y(\theta) = \begin{bmatrix} \cos(\theta/2) & -\sin(\theta/2) \\ \sin(\theta/2) & \cos(\theta/2) \end{bmatrix}$$

$$R_z(\theta) = \begin{bmatrix} e^{-i\theta/2} & 0 \\ 0 & e^{i\theta/2} \end{bmatrix}$$

$$T = \begin{bmatrix} 1 & 0 \\ 0 & e^{-i\pi/4} \end{bmatrix}$$

$$CNOT = \begin{pmatrix} 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 0 & 1 \\ 0 & 0 & 1 & 0 \end{pmatrix}$$

Universal gates, Clifford Gates and non-Clifford gates

5.2 Born Rule

An essential part of quantum computing involves the existence of the Born Rule. Born's measurement rule states that:

$$p(x) = |\langle x | \psi(\theta) \rangle|^2 \quad (5.3)$$

where

$$|\psi(\theta)\rangle = U(\theta)|0\rangle^{\otimes n} \quad (5.4)$$

The state $|\psi(\theta)\rangle$ is generated by evolving state $|0\rangle$ according to a Hamiltonian H that is constructed from gates. Once combined, the gates form a parameterised quantum circuit which is parameterised by using the variables governing each gate, θ . By tuning the values of θ_i one can allow for an evolution to any state that will serve as a solution to a given problem.

By taking the distribution associated to the state, $|\psi(\theta)\rangle$ we can treat the PQC as a generative model, upon measurement will generate samples of a target distribution χ . This model is parameterised by θ , which defines a quantum circuit $U(\theta)$ made up of a set of quantum gates. By measuring the circuit, we can obtain samples. Producing samples that emulate the target distribution involves minimising the parameters of the circuit $U(\theta)$, a process once convergence is reached, will generate accurate samples [25].

5.3 State Preparation

We require state preparation to transfer the classical data onto the Hilbert space. This involves a function ϕ that maps the input vector to an output label. There are many encoding schemes, each of which aim to offer high information density and low error rates; main methods include: basis, amplitude, angle encoding, and QRAM.

Without the use of ancillary qubits, we can expect an exponential circuit depth to prepare an arbitrary quantum state. Using them we can reduce the depth to be sub-exponential scaling, with recent advancements reaching $\Theta(n)$ given $O(n^2)$ ancillary qubits [39, 48].

5.4 Parameterised Quantum Circuits

Parameterised quantum circuits.

6 Quantum Circuit Born Machine

Given a dataset $D = \{x_1, x_2 \dots x_n\}$ consisting of n samples and obeys a given distribution χ_d , we would like the QCBM to learn the distribution and generate synthetic data points that are of the distribution χ_s such that χ_s approximates χ_d . The QCBM is a subclass of parameterised quantum circuits, here the quantum circuit contains parameters which are updated during a training process. The QCBM takes the product state $|0\rangle$ as an input, and through an evolution, transforms into a final state $|\phi_0\rangle$ by a sequence of unitary gates. This can then be measured to obtain a sample of bits $x \sim p_\theta(x_s) = |\langle x | \phi_\theta \rangle|^2$. By training the model we are aiming to let p_θ approach χ_d .

The ansatz for this quantum circuit consists of 7 layers of 1-qubit gates with entangling layers in between them. These are entangled using the CNOT gates as found in the appendix. The number of wires needed depends on the precision required for the generated data. The estimated precision is 12-bit, so the samples are able to take 2^{12} different values in the range of $(v_{min} - \epsilon, v_{max} + \epsilon)$, where $\epsilon > 0$ allows data to be generated that lie outside the range (v_{min}, v_{max}) of the original data.

The QCBM takes a $n \times m$ matrix of parameters in the range $(-\pi, \pi)$ as input, in the form of a dictionary. Each angle takes one of 2^k discrete values, where k is a model parameter. The resulting space therefore spans to: $(2^m)^{n \cdot m}$.

6.1 Barren Plateau

A point of concern when searching for the optimal set of θ s is the large search space, here we may observe issues such as barren plateau(BP). BP insists that the gradient of the parameters of a given PQC will vanish exponentially w.r.t the search space.

Introduce proof of $\frac{\partial C}{\partial \theta} \rightarrow 0$

6.2 Architectures

The design of the ansatz can significantly affect the ability to learn and represent the target distribution. The number of gates, depth, and entanglement structure all affect the expressibility and trainability of the circuit. Circuits with higher entanglement and parameterised gates are theoretically able to represent any distribution, though it comes at the cost of noise and barren plateaus. Choosing a simpler ansatz may converge quicker but results in a weaker approximation, oversimplifying the solution.

The architecture will also affect the optimisation landscape; random parameterised circuits with deep, unstructured layers are more prone to barren plateaus

Due to this we strive for a balance, one that is able to learn the complexities of the market without compromising on finding the optimal parameters.

6.2.1 Brick**6.2.2 Pyramid****6.2.3 Butterfly****6.3 ZX Calculus**

ZX Calculus as a method to optimise/minimise T gate count. introduce what a non clifford gate is.

7 Results

Putting the theory into practice offered insights into the strengths and weaknesses of the model. This section aims to provide quantitative comparisons between the classical and quantum methods. After careful consideration and analysis of parameters for the QCBM, the model used for comparisons against the MJD model is a 13 qubit, 7 layer model. It uses the brick architecture and has been trained for 500 epochs with an Adam optimiser.

7.1 Path Generation

An important part of risk analysis involves path generation, simulating an equity path for the next n days. This provides a range of final values, aiming to simulate price paths accurately in the process. The metrics involved in the analysis involves comparing: the skewness, excess kurtosis, and standard deviation. For a return r_i and mean return \bar{r}

Standard deviation:

$$\sigma = \sqrt{\frac{1}{N} \sum_{i=1}^N (r_i - \bar{r})^2}$$

Skewness:

$$\gamma_1 = \frac{\sum_{i=1}^N (r_i - \bar{r})^3}{(N - 1) \times \sigma^3}$$

Excess kurtosis:

$$\gamma_2 = \frac{1}{N} \frac{\sum_{i=1}^N (r_i - \bar{r})^4}{\sigma^4} - 3$$

These were chosen to highlight the accuracy of the models as well as the ability to represent subtleties in the data such as the asymmetric nature and fat tails that are often present in market data. Both models were first calibrated on Brent Crude Oil stock prices. The test period combines the training and unseen data to see how well the model is able to perform on out-of-sample prices. The period of data includes the recent reaction to Trump's tariffs, April 2025. This was included purposely to observe the model's resilience to tough market activity and heightened volatility.

Parameter	Value
Drift (μ)	0.37956244
Volatility (σ)	0.40736202
Jump Intensity (λ)	530.63931692
Jump Mean (μ_J)	-0.00374471
Jump Variance (σ_J^2)	0.00065703

Table 1: Calibrated Parameters of the MJD Model for Brent Crude Oil

As shown in table 2, the results highlighted that on training data for Brent Crude Oil, the MJD model was able to capture the standard deviation better, however the skewness and excess kurtosis was represented weakly. This is where the QCBM was able to demonstrate superiority, capturing the skewness and excess kurtosis with greater accuracy.

	Original Data	Test Data	MJD Data	QCBM Data
Standard Deviation	0.0233	0.0225	0.02647	0.0268
Skewness	0.6550	0.6709	-0.2579	0.6215
Excess Kurtosis	8.5986	8.8553	7.4809	8.7375

Table 2: Comparison of Brent data with MJD and QCBM Data

	Original Data	Test Data	MJD Data	QCBM Data
Standard Deviation	0.0120	0.0270	0.0126	0.0124
Skewness	0.8602	1.1640	-0.2015	0.7451
Excess Kurtosis	11.5700	5.6518	11.4330	11.3780

Table 3: Comparison of Eurexx Data with MJD and QCBM Data

Table 3 supports most of the arguments made, though has shown slight superiority in representing the fat tails during the training period. As the Eurexx data contains less extreme jumps, we can see the MJD model perform better. That being said, we can still observe the skewness being learnt poorly by the classical model. The test data chosen is a continuation of the training data, both models showing clear difficulties in generalisation. This can be excused however as the metrics differ hugely from the training data. This does raise a point of concern whether these methods employed of market behaviour prediction are appropriate. For this reason, it is more appropriate to treat the models as a synthetic market generator rather than for the purposes of prediction. The differences in the MJD performance should be noted, highlighting how classical methods tend to struggle in tougher market conditions; this is will be further investigated in the following sections.



Figure 2: True price path




Figure 3: Price trajectories for MJD model

A comparison of the log return distributions in figure 4 indicates the lack of tail representation with the MJD model with QCBM's being visible but with higher density. It is important to note that different parameters for the QCBM gave significantly different results. Adding only 1 more qubit led to an over-estimation in kurtosis as well as poor representation of the standard deviation. From figure 5 we can see behaviour that would on first glance mimic the real market. Jumps in the market look realistic, and trajectories take believable paths. The graph of log returns, figure 6 shows us the heightened log returns in the QCBM with occasional peaks as observed in the market data. This raises




Figure 4: Comparison of distributions




Figure 5: Price trajectories for QCBM model

some concerns on whether this would lead to over-hedging if used in a hedging engine. The cumulative sum of returns shows a similarity in the global returns we would expect, the MJD in this regard performs weakly, indicating that the market returned higher returns, an assumption that may lead to under-hedging, leaving an investor exposed to more risk than calculated.




Figure 6: Log returns




Figure 7: Cumulative sum of returns

We can also compare the ACF(Autocorrelation Function) values to see how well the QCBM has replicated the structure of the original dataset. ACF aims to quantify the correlation between observations separated by a lag value, k . This will give us insight into market microstructures within the original data as well as seeing if the QCBM is able to learn these complexities as well. Formally the ACF for a lag k and time series X , can be defined as:

$$\rho_k = \frac{Cov(X_t, X_{t-k})}{\sqrt{Var(X_t) \cdot Var(X_{t-k})}} \quad (7.1)$$

Plots 8 and 9 show the structure of the training data vs the QCBM data. The market data shows significant correlation between the lags, indicating volatility clustering. This is not reflected in the QCBM data, with random spikes showing the model struggles with representing volatility persistence. Similar behaviour is observed in the MJD model as shown in figure 10. We continue this analysis in further detail in the next section.




Figure 8: ACF plot of market data




Figure 9: ACF plot of QCBM data




Figure 10: ACF plot of MJD data

7.2 Volatility Analysis

It is also important to analyse if the local behaviour of the equity paths is well captured. We can do this by analysing the volatility; here the shortfalls of the QCBM become clear. Though the global volatility distribution appears to well represented, using GARCH models shows a disparity in observed volatility clustering compared to the quantum generated volatility. This appears to be a fundamental flaw in the model. The assumption that we can represent a path using [give equation for propagating path] leads to poor volatility clustering as we have treated for each time $0 \leq t_0 \leq \dots \leq t_n$ our random variable $S_{t_r} - S_{t_{r-1}}$ are independent.

First looking at rolling volatility, we observe acceptable performance by the MJD model, displaying realistic market volatility; the QCBM, however, displayed a worse performance, often underestimating volatility peaks, and overestimating noise as shown in figure 11. This heightened noise is also reflected in figure 12 where we see on a weekly basis, the QCBM constantly overestimated the volatility. We can observe the difference in learnt volatility by comparing distributions. Figure 13 shows us how the volatility is skewed with a larger mean and lack of kurtosis. This story worsens as we look at using volatility




Figure 11: Comparison of rolling volatility (20-day window)




Figure 12: Comparison of weekly volatility




Figure 13: Comparison of volatility distributions

models as a comparison method.

GARCH (Generalised Autoregressive Conditional Heteroskedasticity) models were first introduced in 1982 by Robert Eagle [add reference here] to model volatility as a non-constant quantity in financial models. The GARCH(1,1) can be defined as so:

Let r_t be the asset return at time t . This can be decomposed as

$$r_t = \mu + \epsilon_t \quad (7.2)$$

where μ is the mean return and $\epsilon_t = \sigma_t z_t$ where $z \sim N(0,1)$. GARCH models the conditional variance of a given time series process, asset returns in our world, as a function of past squared shocks (ϵ_{t-1}^2) and past conditional variance (σ_{t-1}^2). Using this, our model equation becomes

$$\sigma_t^2 = \omega + \alpha \epsilon_{t-1}^2 + \beta \sigma_{t-1}^2 \quad (7.3)$$

where $\omega > 0$ is the average volatility, $\alpha \geq 0$ is the sensitivity to recent shocks, and $\beta \geq 0$ is the persistence of volatility, the tendency of volatility to be high for periods of time and then low of periods of time.

A further model that improved on GARCH is the Exponential GARCH, another model used for comparison in my findings. This model focussed on asymmetric volatility effects, removing parameter restrictions, and having a logarithmic formulation. We can define an

EGARCH(1,1) as follows

$$\ln(\sigma_t^2) = \omega + \beta \ln(\sigma_{t-1}^2) + \gamma \frac{\epsilon_{t-1}}{\sigma_{t-1}} + \alpha \left(\frac{|\epsilon_{t-1}|}{\sigma_{t-1}} - \sqrt{\frac{2}{\pi}} \right) \quad (7.4)$$

where the extra term γ accounts for the leverage effect; the negative correlation between asset returns and volatility change often observed in markets.




Figure 14: EGARCH model fit to models

The implementation and parameter estimation for ω, α, β and γ were all handled by the python package 'arch' [insert this reference <https://pypi.org/project/arch/>]. Calibrating a EGARCH(1,1) to the different model returns gave the following graph. In figure 14 we can observe the MJD model displaying more accurate clustering compared to the QCBM which remains conservative. This may suggest that the MJD model is more suitable for predicting volatility over a short period of time. As an improvement to the QCBM, it may be of use to create a hybrid model, one that combines the QCBM and a volatility model of choice.

7.3 VaR & CVaR

To determine the usefulness of the generator for hedging purposes, we must explore how both generators perform in the extreme percentiles. These are the scenarios that we call Black-Swan events; an event that is unexpected, infrequent but has huge financial implications. Not being hedged against such changes can leave an investor in unwanted positions, exposed to large changes and at risk of losing a lot of money. To quantify expected loss a market may give, we employ techniques such as VaR(Value at Risk) and CVaR(Conditional Value at Risk).

VaR is a measure that focuses on the loss at a given percentile under normal market conditions. It provides a threshold such that the probability of a loss exceeding a given value is a chosen percentile i.e 1% or 5%. An intuitive way to think about it is 1% VaR of -0.05 means there is a 1% chance of losing at least 5% of the portfolio value. Mathematically we can define VaR:

Let X be a random variable with cdf $F_X(z) = P\{X \leq z\}$, this is our loss function. The VaR of X at a confidence level α is given by

$$VaR_\alpha(X) = \min\{z | F_X(z) \geq \alpha\} \quad (7.5)$$

VaR, however, has a few crucial limitations. VaR does not say anything about the size of losses made after the given α value. This makes it very hard to quantify the loss an investor may be exposed to. VaR also assumes liquidity in the market at any position when in reality during periods of market stress, the ability to buy or sell at any price can be very difficult. It should also be known that VaR has a non-subadditivity constraint, meaning VaR for different portfolios cannot be added together. For these reasons we opt to use Conditional Value at Risk, a more sophisticated tool designed to overcome the given limitations.

CVaR focusses on measuring the risk of extreme losses. Instead of considering of considering the probability of losing a given amount, we instead focus on how much is lost in the given quantile. This quantifies our tail risk, making it appropriate for evaluating the heavy tails of our market generators. Intuitively a 1% CVaR of -0.05 means that, in the worst 1% of cases, the average loss is 5%. We can define CVaR at a level α as

$$CVaR_\alpha(X) = \mathbb{E}[X | X \geq VaR_\alpha(X)] \quad (7.6)$$

This provides us with a sturdier framework to evaluate performances although we must still be careful of its assumptions and limitations. We must assume that the distribution of the loss is measured accurately and market conditions are also represented faithfully within the given time window. Once we accept these to be true, we are provided with a method to compare strategies and generators, as well as conforming to financial regulations such as Basel III.

The first test was on Eurexx data from the COVID period. This was excluded from the training dataset to be used as an extreme market situation. Market generators need to be resilient and able to account such market shocks, these are the scenarios where hedging

	VaR (1%)	CVaR (1%)
Original Data	-0.0171	-0.0174
QCBM Simulated	-0.0291	-0.0308
MJDM Simulated	-0.0360	-0.0377

Table 4: VaR & CVaR at the 1% Quantile for COVID data

	VaR (5%)	CVaR (5%)
Original Data	-0.0159	-0.0167
QCBM Simulated	-0.0223	-0.0270
MJDM Simulated	-0.0283	-0.0338

Table 5: VaR & CVaR at the 5% Quantile for COVID data

positions stop investors from losing large amounts of money. At the 1% quantile we can see both the MJD and QCBM exhibit greater losses in the tails compared to the real data. The QCBM has a VaR of 2.91% and CVaR of 3.08%, these are both significant overestimates of tail risk present in the training data. The MJD overestimates the tail risk even greater, expecting much larger than expected losses. We could explain this due to the jump component. These pessimistic scenarios, though stopping large losses, in reality may lead to over hedging, a situation where a given investor may be too protected to the market, limiting potential gains. These scenarios may be suboptimal for a deep hedging engine, where the optimal hedge is going to be larger than needed.

At the 5% quantile we see a similar story fold out, both generators estimating the tail risk to be greater than observed. The QCBM once again provides a closer score to the data, indicating the model is more suitable for hedging purposes, though only marginally. What should be noted is the large gap between the VaR and CVaR, indicating the MJD model has accounted for significant market shocks, indicating high sensitivity to extreme events.

The next scenario is on the unseen continuation of the Brent Crude Oil dataset, though not excluded purposely than just for out-of-sample testing, the market conditions observed have been very volatile. As mentioned earlier, the effect of Trump's tariffs were felt in all markets, particularly affecting assets denominated in the US dollar. Though as unfortunate as the scenario may be to investors, this has provided me with an excellent test for the two models.

	VaR (1%)	CVaR (1%)
Original Data	-0.0300	-0.0342
QCBM Simulated	-0.0429	-0.0524
MJDM Simulated	-0.0522	-0.0631

Table 6: VaR & CVaR at the 1% Quantile for Brent Crude Data

	VaR (5%)	CVaR (5%)
Original Data	-0.0239	-0.0274
QCBM Simulated	-0.0313	-0.0391
MJDM Simulated	-0.0414	-0.0508

Table 7: VaR & CVaR at the 5% Quantile for Brent Crude Data

Though the data was more volatile, both models demonstrated that they are able to capture the tail risk. The training dataset had large amounts of volatility so we would expect good representation of tail risk. As seen with the COVID results, the models are overestimating the downside risk with scenarios showing greater loss than expected. The large gap between the VaR and CVaR values once again tells us that the underlying models predict tail events far higher than observed in the market. It should be said that the QCBM model is able to reflect the tail risk more accurately than the MJD, with values that are closer to the original data. One might make conclusions that this makes the QCBM a useful model but the poor performance to the original data raises concerns; we can however say that the model is more appropriate to use than the MJD in this scenario.

7.4 Hedging

7.5 Barren Plateau

7.5.1 ZX Calculus

8 Outlook and Conclusions

[add in the conclusion, large hedges are suboptimal, and traders underhedge anyways]

9.1 QCBM Architectures



References

- [1] John Armstrong and George Tatlow. *Deep Gamma Hedging*. Sept. 20, 2024. DOI: 10.48550/arXiv.2409.13567. arXiv: 2409.13567[q-fin]. URL: <http://arxiv.org/abs/2409.13567> (visited on 12/04/2024).
- [2] Marcello Benedetti et al. “A generative modeling approach for benchmarking and training shallow quantum circuits”. In: *npj Quantum Information* 5.1 (May 27, 2019), p. 45. ISSN: 2056-6387. DOI: 10.1038/s41534-019-0157-8. URL: <https://www.nature.com/articles/s41534-019-0157-8> (visited on 11/21/2024).
- [3] Marcello Benedetti et al. “Erratum: Parameterized quantum circuits as machine learning models (2019 *Quant. Sci. Tech.* 4 043001)”. In: *Quantum Science and Technology* 5.1 (Dec. 4, 2019), p. 019601. ISSN: 2058-9565. DOI: 10.1088/2058-9565/ab5944. URL: <https://iopscience.iop.org/article/10.1088/2058-9565/ab5944> (visited on 11/21/2024).
- [4] Michael R. Berthold, Ad Feelders, and Georg Kreml, eds. *Advances in Intelligent Data Analysis XVIII: 18th International Symposium on Intelligent Data Analysis, IDA 2020, Konstanz, Germany, April 27–29, 2020, Proceedings*. Vol. 12080. Lecture Notes in Computer Science. Cham: Springer International Publishing, 2020. ISBN: 978-3-030-44583-6 978-3-030-44584-3. DOI: 10.1007/978-3-030-44584-3. URL: <http://link.springer.com/10.1007/978-3-030-44584-3> (visited on 12/04/2024).
- [5] Bloomberg. *Bloomberg*. Bloomberg.com, 2023. URL: <https://www.bloomberg.com/>.
- [6] Nicolas Boursin, Carl Remlinger, and Joseph Mikael. “Deep Generators on Commodity Markets Application to Deep Hedging”. In: *Risks* 11.1 (Dec. 23, 2022), p. 7. ISSN: 2227-9091. DOI: 10.3390/risks11010007. URL: <https://www.mdpi.com/2227-9091/11/1/7> (visited on 11/21/2024).
- [7] Hans Buehler, Murray Phillip, and Ben Wood. “Deep Bellman Hedging”. In: *SSRN Electronic Journal* (2022). ISSN: 1556-5068. DOI: 10.2139/ssrn.4151026. URL: <https://www.ssrn.com/abstract=4151026> (visited on 11/21/2024).
- [8] Hans Buehler et al. “Deep Hedging: Hedging Derivatives Under Generic Market Frictions Using Reinforcement Learning”. In: *SSRN Electronic Journal* (2019). ISSN: 1556-5068. DOI: 10.2139/ssrn.3355706. URL: <https://www.ssrn.com/abstract=3355706> (visited on 11/21/2024).
- [9] El Amine Cherrat et al. “Quantum Deep Hedging”. In: *Quantum* 7 (Nov. 29, 2023), p. 1191. ISSN: 2521-327X. DOI: 10.22331/q-2023-11-29-1191. arXiv: 2303.16585[quant-ph]. URL: <http://arxiv.org/abs/2303.16585> (visited on 12/04/2024).
- [10] Vedant Choudhary, Sebastian Jaimungal, and Maxime Bergeron. *FuNVol: A Multi-Asset Implied Volatility Market Simulator using Functional Principal Components and Neural SDEs*. Dec. 26, 2023. DOI: 10.48550/arXiv.2303.00859. arXiv: 2303.00859[q-fin]. URL: <http://arxiv.org/abs/2303.00859> (visited on 12/04/2024).

- [11] Eliahu Cohen. “What Weak Measurements and Weak Values Really Mean: Reply to Kastner”. In: *Foundations of Physics* 47.10 (Oct. 2017), pp. 1261–1266. ISSN: 0015-9018, 1572-9516. DOI: 10.1007/s10701-017-0107-2. URL: <http://link.springer.com/10.1007/s10701-017-0107-2> (visited on 11/21/2024).
- [12] Kalyan Dasgupta and Binoy Paine. *Loading Probability Distributions in a Quantum circuit*. Aug. 29, 2022. arXiv: 2208.13372[quant-ph]. URL: <http://arxiv.org/abs/2208.13372> (visited on 11/21/2024).
- [13] Julien Dudas et al. “Quantum reservoir computing implementation on coherently coupled quantum oscillators”. In: *npj Quantum Information* 9.1 (July 7, 2023), p. 64. ISSN: 2056-6387. DOI: 10.1038/s41534-023-00734-4. URL: <https://www.nature.com/articles/s41534-023-00734-4> (visited on 11/21/2024).
- [14] *EURO STOXX 50*. Wikipedia, Aug. 2021. URL: https://en.wikipedia.org/wiki/EURO_STOXX_50.
- [15] Simon Fecamp, Joseph Mikael, and Xavier Warin. “Deep learning for discrete-time hedging in incomplete markets”. In: *The Journal of Computational Finance* (2021). ISSN: 14601559, 17552850. DOI: 10.21314/JCF.2021.006. URL: <https://www.risk.net/journal-of-computational-finance/7871526/deep-learning-for-discrete-time-hedging-in-incomplete-markets> (visited on 12/04/2024).
- [16] Giacomo Franceschetto et al. *Harnessing quantum back-action for time-series processing*. Nov. 6, 2024. arXiv: 2411.03979[quant-ph]. URL: <http://arxiv.org/abs/2411.03979> (visited on 11/21/2024).
- [17] Keisuke Fujii and Kohei Nakajima. *Quantum reservoir computing: a reservoir approach toward quantum machine learning on near-term quantum devices*. Nov. 10, 2020. arXiv: 2011.04890[quant-ph]. URL: <http://arxiv.org/abs/2011.04890> (visited on 11/21/2024).
- [18] Santanu Ganguly. “Implementing Quantum Generative Adversarial Network (qGAN) and QCBM in Finance”. In: ().
- [19] Jorge García-Bení et al. “Squeezing as a resource for time series processing in quantum reservoir computing”. In: *Optics Express* 32.4 (Feb. 12, 2024), p. 6733. ISSN: 1094-4087. DOI: 10.1364/OE.507684. arXiv: 2310.07406[quant-ph]. URL: <http://arxiv.org/abs/2310.07406> (visited on 11/21/2024).
- [20] Kaitlin Gili et al. *Do Quantum Circuit Born Machines Generalize?* arXiv.org, 2022. URL: <https://arxiv.org/abs/2207.13645>.
- [21] L. C. G. Govia et al. “Quantum reservoir computing with a single nonlinear oscillator”. In: *Physical Review Research* 3.1 (Jan. 25, 2021), p. 013077. ISSN: 2643-1564. DOI: 10.1103/PhysRevResearch.3.013077. URL: <https://link.aps.org/doi/10.1103/PhysRevResearch.3.013077> (visited on 11/21/2024).
- [22] Haim Horowitz, Pooja Rao, and Santosh Kumar Radha. *A quantum generative model for multi-dimensional time series using Hamiltonian learning*. Apr. 13, 2022. DOI: 10.48550/arXiv.2204.06150. arXiv: 2204.06150[quant-ph]. URL: <http://arxiv.org/abs/2204.06150> (visited on 12/04/2024).

- [23] Iordanis Kerenidis, Jonas Landman, and Natansh Mathur. *Classical and Quantum Algorithms for Orthogonal Neural Networks*. Dec. 23, 2022. arXiv: 2106.07198[quant-ph]. URL: <http://arxiv.org/abs/2106.07198> (visited on 11/21/2024).
- [24] Alex Kondratyev. “Non-Differentiable Learning of Quantum Circuit Born Machine with Genetic Algorithm”. In: *Wilmott* 2021.114 (July 2021), pp. 50–61. ISSN: 1540-6962, 1541-8286. DOI: 10.1002/wilm.10943. URL: <https://onlinelibrary.wiley.com/doi/10.1002/wilm.10943> (visited on 11/21/2024).
- [25] Jin-Guo Liu and Lei Wang. “Differentiable Learning of Quantum Circuit Born Machine”. In: *Physical Review A* 98.6 (Dec. 19, 2018), p. 062324. ISSN: 2469-9926, 2469-9934. DOI: 10.1103/PhysRevA.98.062324. arXiv: 1804.04168[quant-ph]. URL: <http://arxiv.org/abs/1804.04168> (visited on 11/21/2024).
- [26] *LSEG Workspace*. www.lseg.com. URL: <https://www.lseg.com/en/data-analytics/products/workspace>.
- [27] Ali Asghar Movahed and Houshyar Noshad. “Introducing a new approach for modeling a given time series based on attributing any random variation to a jump event: jump-jump modeling”. In: *Scientific Reports* 14.1 (Jan. 12, 2024), p. 1234. ISSN: 2045-2322. DOI: 10.1038/s41598-024-51863-5. URL: <https://www.nature.com/articles/s41598-024-51863-5> (visited on 11/21/2024).
- [28] Pere Mujal et al. “Time-series quantum reservoir computing with weak and projective measurements”. In: *npj Quantum Information* 9.1 (Feb. 23, 2023), p. 16. ISSN: 2056-6387. DOI: 10.1038/s41534-023-00682-z. URL: <https://www.nature.com/articles/s41534-023-00682-z> (visited on 11/21/2024).
- [29] Bernt Øksendal. “Application to Mathematical Finance”. In: *Stochastic Differential Equations: An Introduction with Applications*. Berlin, Heidelberg: Springer Berlin Heidelberg, 2003, pp. 269–313. ISBN: 978-3-642-14394-6. DOI: 10.1007/978-3-642-14394-6_12. URL: https://doi.org/10.1007/978-3-642-14394-6_12.
- [30] Anupama Padha and Anita Sahoo. “Quantum deep neural networks for time series analysis”. In: *Quantum Information Processing* 23.6 (May 24, 2024), p. 205. ISSN: 1573-1332. DOI: 10.1007/s11128-024-04404-y. URL: <https://link.springer.com/10.1007/s11128-024-04404-y> (visited on 11/21/2024).
- [31] Annie E. Paine, Vincent E. Elfving, and Oleksandr Kyriienko. “Quantum Quantile Mechanics: Solving Stochastic Differential Equations for Generating Time-Series”. In: *Advanced Quantum Technologies* 6.10 (Oct. 2023), p. 2300065. ISSN: 2511-9044, 2511-9044. DOI: 10.1002/qute.202300065. arXiv: 2108.03190[quant-ph]. URL: <http://arxiv.org/abs/2108.03190> (visited on 11/21/2024).
- [32] Obdulia Pichardo Lagunas, Juan Martínez-Miranda, and Bella Martínez Seis, eds. *Advances in Computational Intelligence: 21st Mexican International Conference on Artificial Intelligence, MICAI 2022, Monterrey, Mexico, October 24–29, 2022, Proceedings, Part I*. Vol. 13612. Lecture Notes in Computer Science. Cham: Springer Nature Switzerland, 2022. ISBN: 978-3-031-19492-4 978-3-031-19493-1. DOI: 10.1007/978-3-031-19493-1. URL: <https://link.springer.com/10.1007/978-3-031-19493-1> (visited on 11/21/2024).

- [33] Chunhui Qiao and Xiangwei Wan. *Enhancing Black-Scholes Delta Hedging via Deep Learning*. Aug. 24, 2024. DOI: 10.48550/arXiv.2407.19367. arXiv: 2407.19367[q-fin]. URL: <http://arxiv.org/abs/2407.19367> (visited on 12/04/2024).
- [34] Lupei Qin, Wei Feng, and Xin-Qi Li. “Simple understanding of quantum weak values”. In: *Scientific Reports* 6.1 (Feb. 3, 2016), p. 20286. ISSN: 2045-2322. DOI: 10.1038/srep20286. URL: <https://www.nature.com/articles/srep20286> (visited on 11/21/2024).
- [35] Michael Ragone et al. “A Lie algebraic theory of barren plateaus for deep parameterized quantum circuits”. In: *Nature Communications* 15 (Aug. 2024). DOI: 10.1038/s41467-024-49909-3. (Visited on 11/19/2024).
- [36] Samir Saissi Hassani and Georges Dionne. “The New International Regulation of Market Risk: Roles of VaR and CVaR in Model Validation”. In: *SSRN Electronic Journal* (2021). ISSN: 1556-5068. DOI: 10.2139/ssrn.3766511. URL: <https://www.ssrn.com/abstract=3766511> (visited on 12/04/2024).
- [37] Samir Saissi Hassani and Georges Dionne. “The New International Regulation of Market Risk: Roles of VaR and CVaR in Model Validation”. In: *SSRN Electronic Journal* (2021). DOI: 10.2139/ssrn.3766511. (Visited on 01/16/2022).
- [38] Martin Schweizer and Johannes Wissel. “Arbitrage-free market models for option prices: the multi-strike case”. In: *Finance and Stochastics* 12.4 (Oct. 2008), pp. 469–505. ISSN: 0949-2984, 1432-1122. DOI: 10.1007/s00780-008-0068-6. URL: <http://link.springer.com/10.1007/s00780-008-0068-6> (visited on 12/04/2024).
- [39] Ali Shaib et al. “Efficient noise mitigation technique for quantum computing”. In: *Scientific Reports* 13.1 (Mar. 8, 2023), p. 3912. ISSN: 2045-2322. DOI: 10.1038/s41598-023-30510-5. URL: <https://www.nature.com/articles/s41598-023-30510-5> (visited on 11/21/2024).
- [40] Lina Song. “Dynamic Modeling and Simulation of Option Pricing Based on Fractional Diffusion Equations with Double Derivatives”. In: *Computational Economics* (May 26, 2024). ISSN: 0927-7099, 1572-9974. DOI: 10.1007/s10614-024-10628-y. URL: <https://link.springer.com/10.1007/s10614-024-10628-y> (visited on 12/04/2024).
- [41] Apimuk Sornsaeng, Ninnat Dangniam, and Thiparat Chotibut. “Quantum next generation reservoir computing: an efficient quantum algorithm for forecasting quantum dynamics”. In: *Quantum Machine Intelligence* 6.2 (Dec. 2024), p. 57. ISSN: 2524-4906, 2524-4914. DOI: 10.1007/s42484-024-00188-7. URL: <https://link.springer.com/10.1007/s42484-024-00188-7> (visited on 11/21/2024).
- [42] Nikitas Stamatopoulos and William J. Zeng. “Derivative Pricing using Quantum Signal Processing”. In: *Quantum* 8 (Apr. 30, 2024), p. 1322. ISSN: 2521-327X. DOI: 10.22331/q-2024-04-30-1322. arXiv: 2307.14310[quant-ph]. URL: <http://arxiv.org/abs/2307.14310> (visited on 11/21/2024).

- [43] György Steinbrecher and William T. Shaw. “Quantile mechanics”. In: *European Journal of Applied Mathematics* 19.2 (Apr. 2008). ISSN: 0956-7925, 1469-4425. DOI: 10.1017/S0956792508007341. URL: http://www.journals.cambridge.org/abstract_S0956792508007341 (visited on 12/04/2024).
- [44] Andrew Vlasic. *Quantum Circuits, Feature Maps, and Expanded Pseudo-Entropy: A Categorical Theoretic Analysis of Encoding Real-World Data into a Quantum Computer*. Oct. 29, 2024. DOI: 10.48550/arXiv.2410.22084. arXiv: 2410.22084[quant-ph]. URL: <http://arxiv.org/abs/2410.22084> (visited on 12/04/2024).
- [45] Magnus Wiese et al. “Deep Hedging: Learning to Simulate Equity Option Markets”. In: *SSRN Electronic Journal* (2019). ISSN: 1556-5068. DOI: 10.2139/ssrn.3470756. URL: <https://www.ssrn.com/abstract=3470756> (visited on 11/21/2024).
- [46] Magnus Wiese et al. *Multi-Asset Spot and Option Market Simulation*. Dec. 13, 2021. DOI: 10.48550/arXiv.2112.06823. arXiv: 2112.06823[q-fin]. URL: <http://arxiv.org/abs/2112.06823> (visited on 12/04/2024).
- [47] Neil A. Wilmot and Charles F. Mason. “Jump Processes in the Market for Crude Oil”. In: *The Energy Journal* 34.1 (Jan. 2013), pp. 33–48. ISSN: 0195-6574, 1944-9089. DOI: 10.5547/01956574.34.1.2. URL: <https://journals.sagepub.com/doi/10.5547/01956574.34.1.2> (visited on 11/21/2024).
- [48] Xiao-Ming Zhang, Tongyang Li, and Xiao Yuan. “Quantum State Preparation with Optimal Circuit Depth: Implementations and Applications”. In: *Physical Review Letters* 129.23 (Nov. 30, 2022), p. 230504. ISSN: 0031-9007, 1079-7114. DOI: 10.1103/PhysRevLett.129.230504. arXiv: 2201.11495[quant-ph]. URL: <http://arxiv.org/abs/2201.11495> (visited on 11/21/2024).
- [49] Christa Zoufal, Aurélien Lucchi, and Stefan Woerner. “Quantum Generative Adversarial Networks for learning and loading random distributions”. In: *npj Quantum Information* 5.1 (Nov. 22, 2019), p. 103. ISSN: 2056-6387. DOI: 10.1038/s41534-019-0223-2. URL: <https://www.nature.com/articles/s41534-019-0223-2> (visited on 11/21/2024).